

# The PPC64 Port of VirtualThreads

Richard Reingruber, SAP  
October, 2022

# Agenda

VirtualThreads Refresher

VM Continuations

- Abstract Operations
- Implementation based on StackChunks

”Freezing” / Reallocating Frames from Stack into StackChunks on Heap

“Thawing” vice versa

GC Integration

Potential Performance Cliffs

Integration with established internal APIs related to Threads and Stacks

# Virtual Threads (aka Project Loom)

Lightweight user-mode threads hiding non-blocking IO

Allow for more concurrency for higher throughput

Comparatively Simple Programming Model

- facilitates Development/Maintenance

Compatible

- With existing thread APIs
- With existing tools
- Beware limitations, most notably: active synchronized blocks/methods prevent lightweight context switching

Virtual threads can significantly improve application throughput

- If the number of concurrent tasks is high (more than a few thousands)
- If the workload is not CPU-bound

<https://openjdk.org/jeps/425>

# VM Continuations for Switching Virtual Threads (VT)

VM Support required (-XX:+VMContinuations)

A (VM) **Continuation** is a Java Object C

- Instantiated with a lambda expression  $\lambda$  or Runnable (just like a thread)
- $\lambda$  can be executed by calling the **run()** method of C; run() blocks while C executes
- C can **self-suspend** execution by calling its **yield()** method.  
run() will then unblock and return to its caller while yield() blocks
- C can be **resumed** by calling the **run()** method of C again.  
The yield() from the suspend above will then return and execution continues.

Starting a VT means creating a Continuation and running it

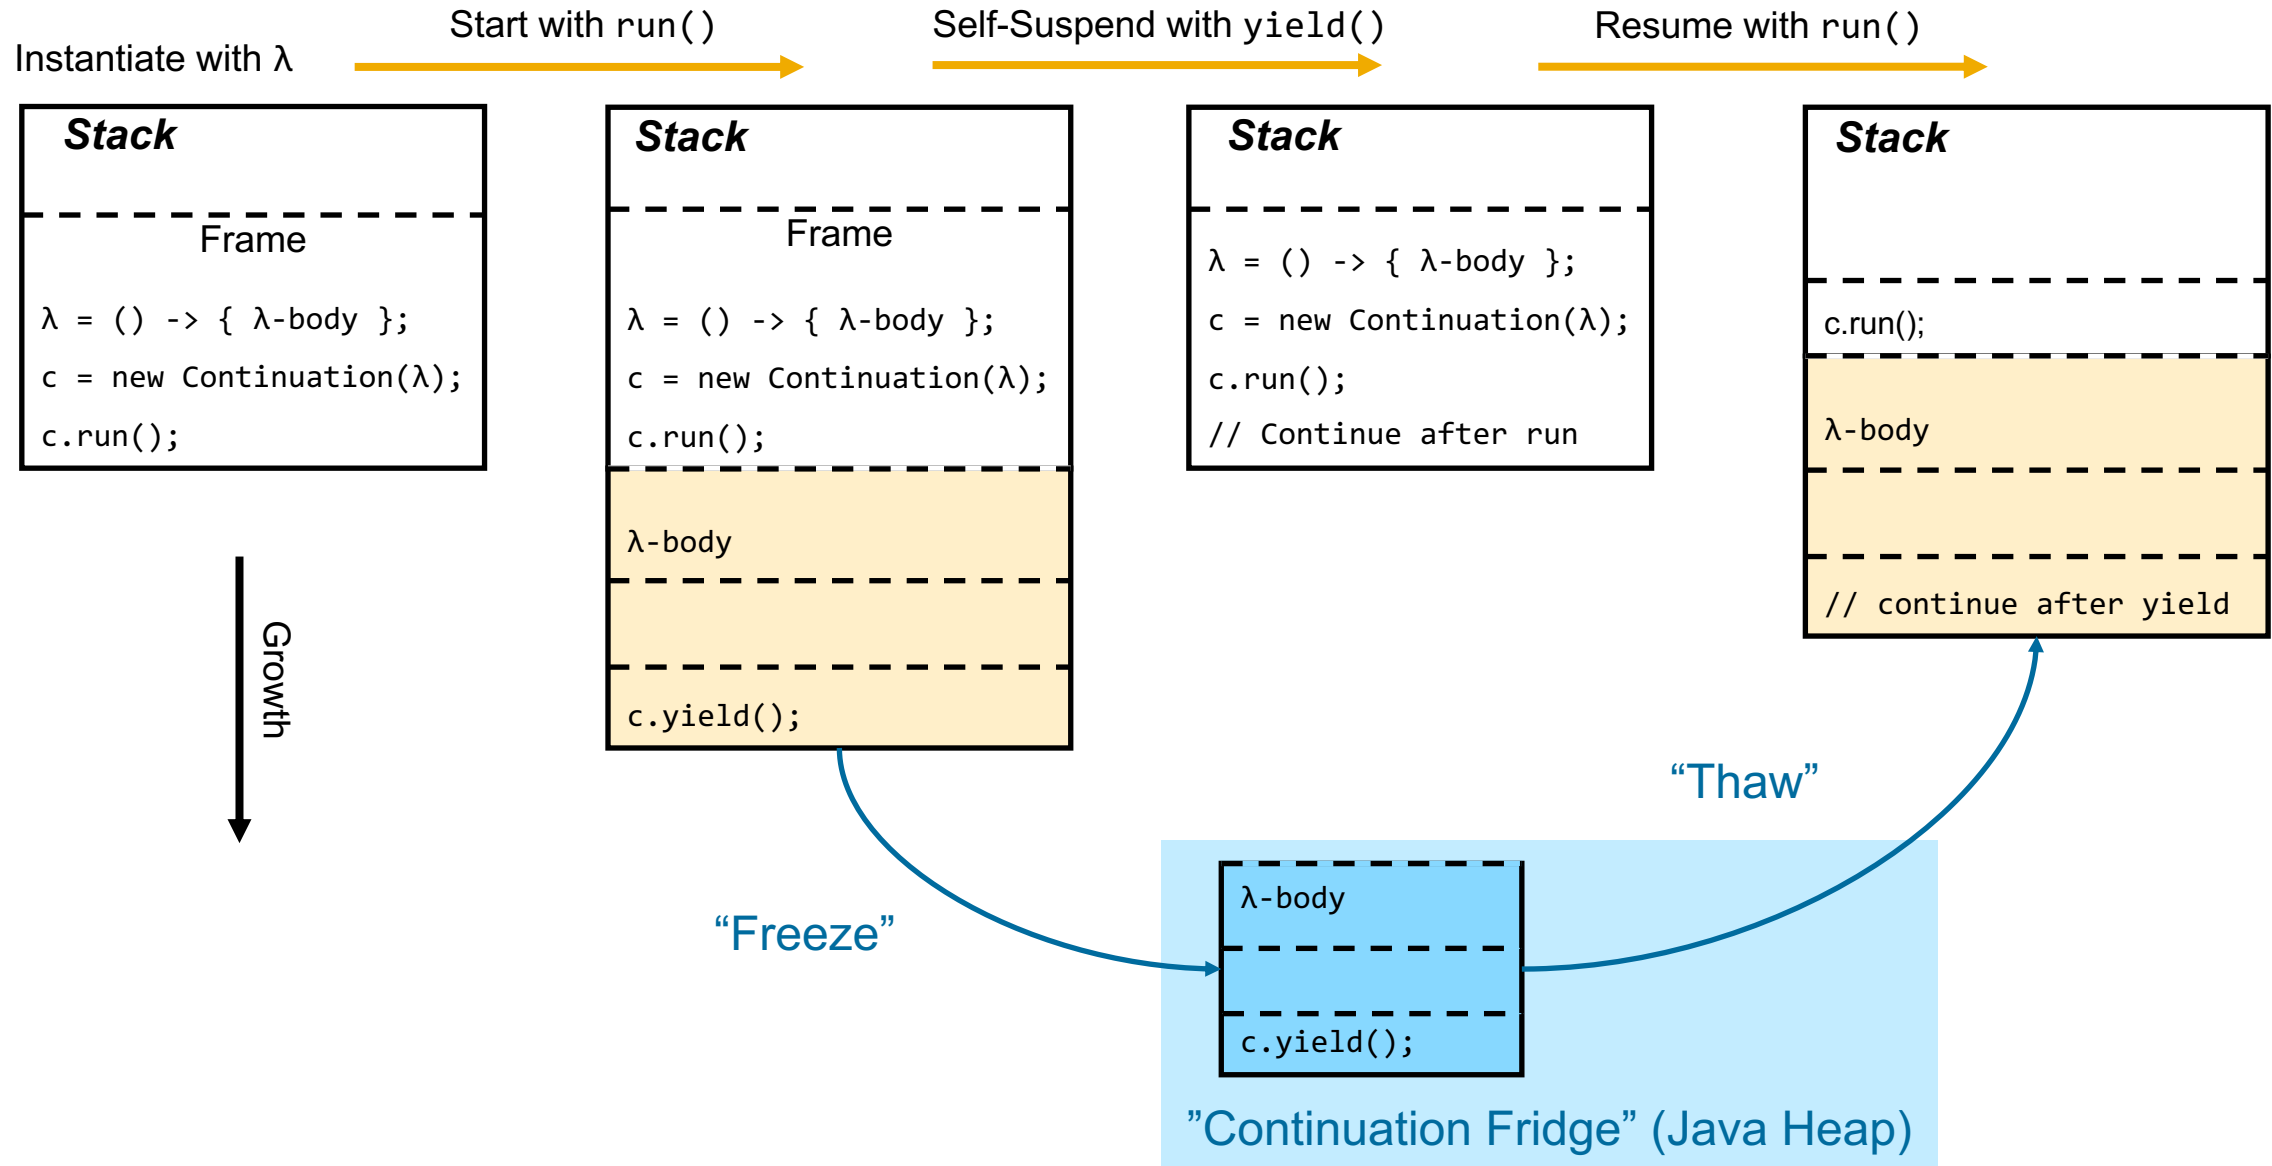
The VT Scheduler makes use of the **resume** and **suspend** operations to switch VTs

Threads are not preempted depending on time  
but only in classlib calls that block (synchronization, IO)

Continuations are private jdk internal objects.

Without VM support (-XX:-VMContinuations) VT are implemented as std. heavy-weight threads

# Visualization of Continuation Operations




# Continuations and StackChunks

`Continuation` and `StackChunk` are classes in the package `jdk.internal.vm`

A Continuation has a list of StackChunks

The StackChunks store the “frozen” frames of a Continuation on heap

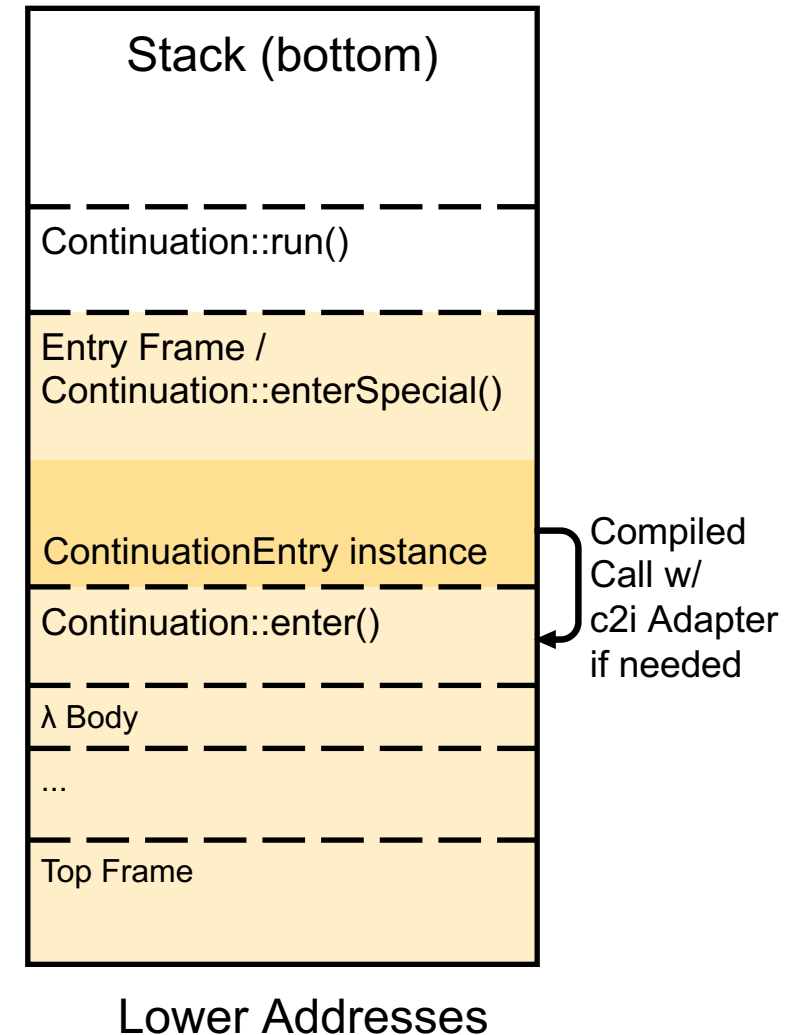
- Top frame has lowest address (just like on the stack)
- StackChunks are of variable length and require special handling during GC
- New Klass type to manage this: `InstanceStackChunkKlass`



# enterSpecial Intrinsic, Entry Frame, ContinuationEntry

`j.i.vm.Continuation::enterSpecial()`

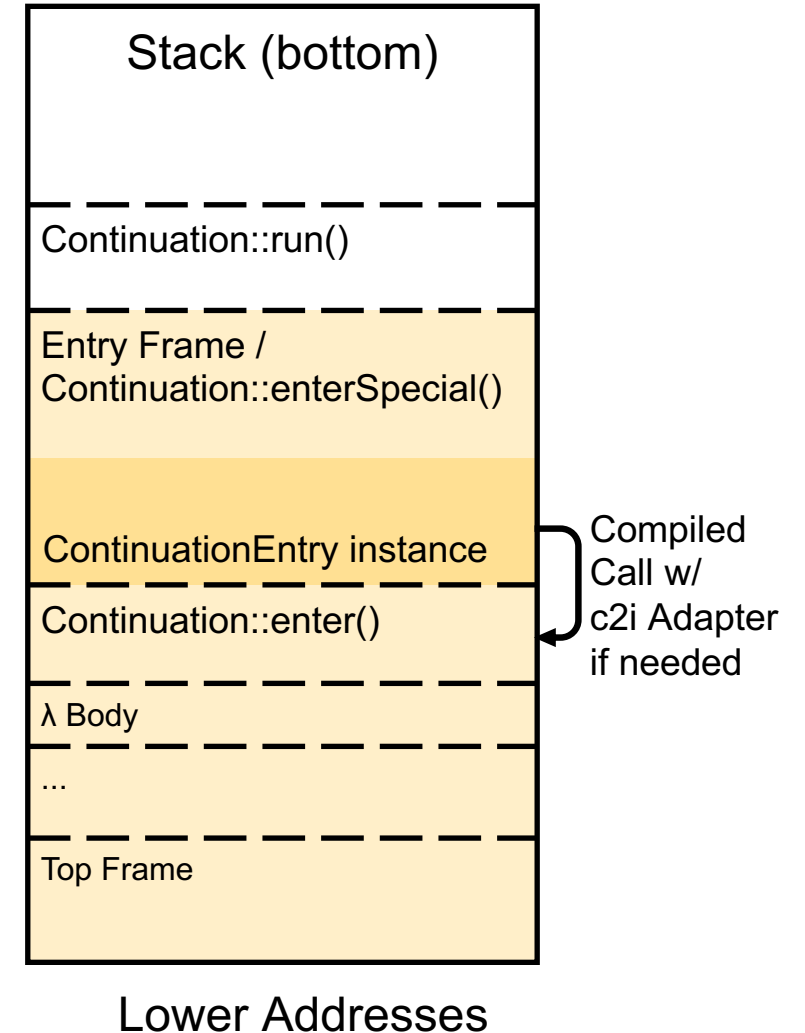
- declared as native method, implemented as a vm intrinsic
- Special native wrapper
  - nmethod that is called with the compiled (jit) calling convention
  - handwritten generated assembler
- Pushes the **Entry Frame** with the **ContinuationEntry** instance, initializes it, and prepends it to the `JavaThread`'s list of CEs
- Depending on parameter `isContinue` either
  - `false`: a static compiled call is executed. There is special handling in the runtime that binds the call to `Continuation::enter()` which calls the continuations  $\lambda$
  - `true`: stack frames are copied ("thawed") from the heap to the stack and control is transferred to the top frame
- See [gen\\_continuation\\_enter\(\)](#)



# enterSpecial Intrinsic, Entry Frame, ContinuationEntry

## ContinuationEntry instance

- located in the Entry Frame
- `JavaThread` has a list of its `ContinuationEntry` instances
- has a Reference to `j.i.vm.Continuation` instance
- Other continuation related data copied from the `JavaThread` for the parent continuation (e.g. locked monitor count)

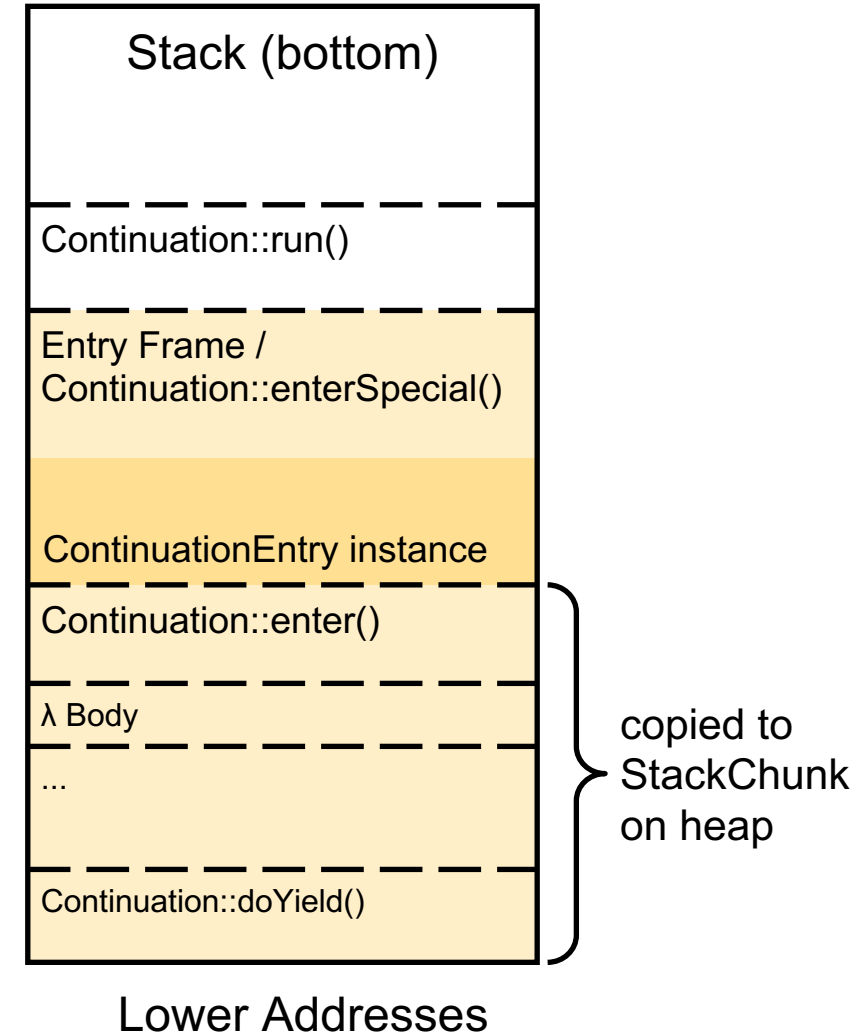




# doYield Intrinsic

`j.i.vm.Continuation::doYield()`

- declared as native method, implemented as a vm intrinsic
- Special native wrapper
  - nmethod that is called with the compiled (jit) calling convention
  - handwritten generated assembler
- Calls the runtime where the frames atop the entry frame are copied (“frozen”) to a StackChunk instance on the heap
- Pops frozen frames by restoring the ContinuationEntry\* from the JavaThread.
- Copies the parent’s data in the ContinuationEntry back to the JavaThread and removes this CE as list head
- Pops the entry frame and returns control to its caller
- See [gen\\_continuation\\_yield\(\)](#)



# Freezing Frames – Slow Path

Fail if `JavaThread::_held_monitor_count > 0`

Iterate frames recursively top to bottom

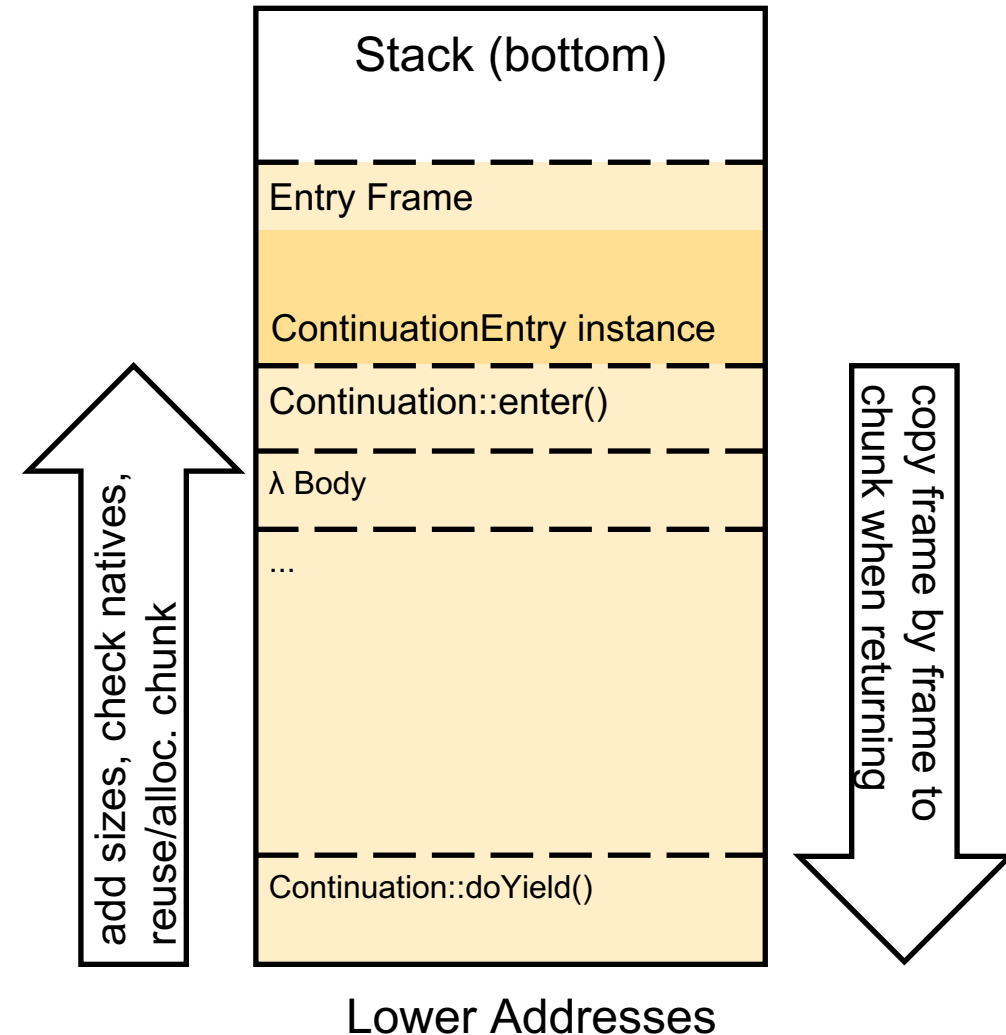
- adding the sizes
- fail if native frame found

At the bottom frame

- required size *S* is known
- existing `StackChunk` can be reused
  - if  $S \leq \text{free in}$
  - if never traversed by GC
  - if it does not need special GC barriers (Shenandoah, ZGC)
- Otherwise a new `StackChunk` is allocated

Then Copy, frame by frame, bottom to top to `StackChunk`

Code: `FreezeBase::new_heap_frame()` and callers




# Freezing Frames – Slow Path – Sizing

## Frame Size

- straight forward: decrement of the stack pointer (SP) before the body of a function is executed
- Beware details, e.g. interpreter frames can grow when stack locks are allocated


## Size function $\Sigma$

- size occupied by a frame (part) in the StackChunk
- Frame: size of complete state!
  - This includes args passed on stack
  - Even though stack args reside in the caller

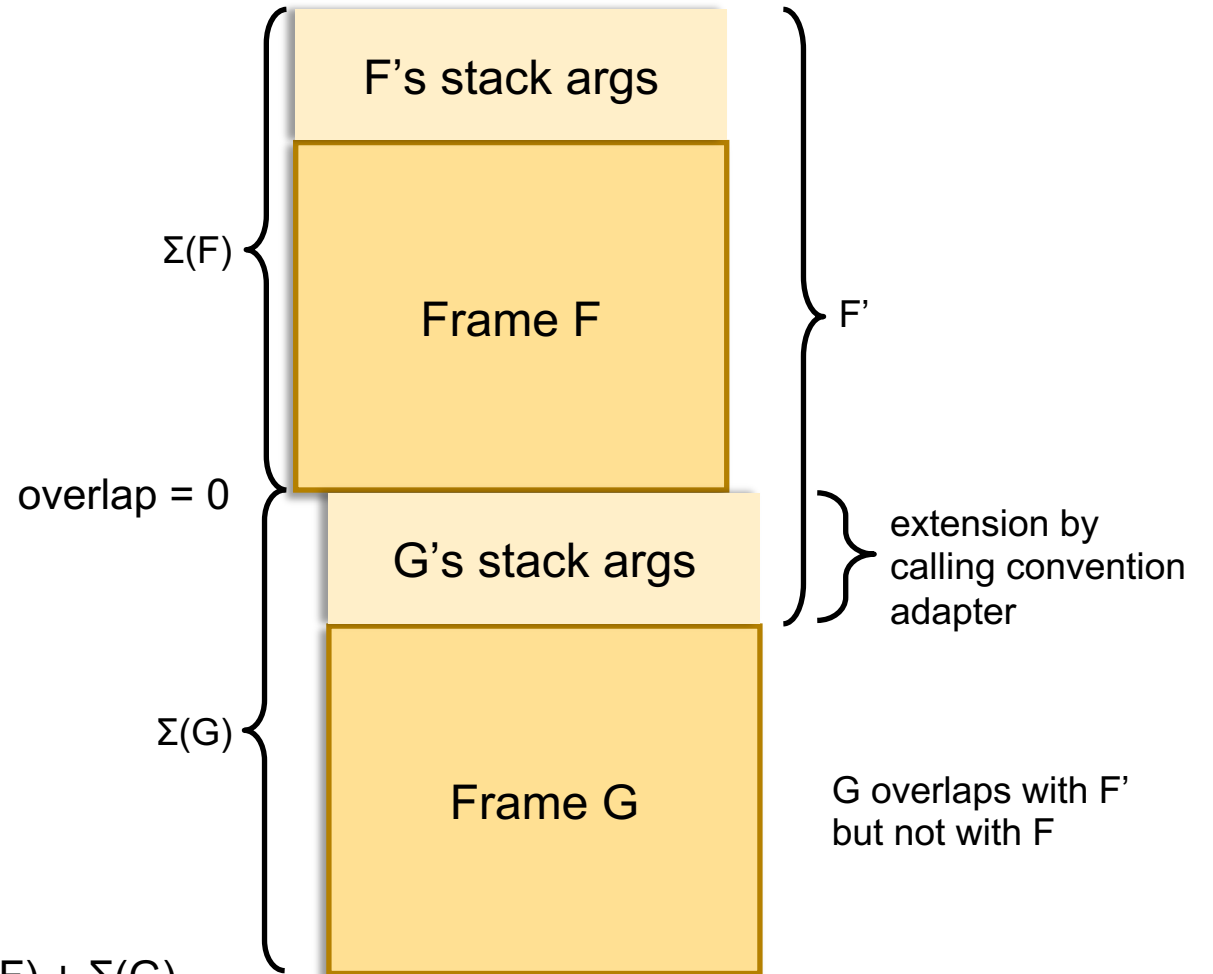


# Freezing Frames – Slow Path – Sizing Caller/Callee

Same Frame Type: stackargs in Caller Frame



Different Frame Type: stackargs in ext. Caller Frame



$$\Sigma(F, G) = \Sigma(F) + \Sigma(G) - \text{overlap}$$

## Freezing Frames – Slow Path – Sizing with Metadata

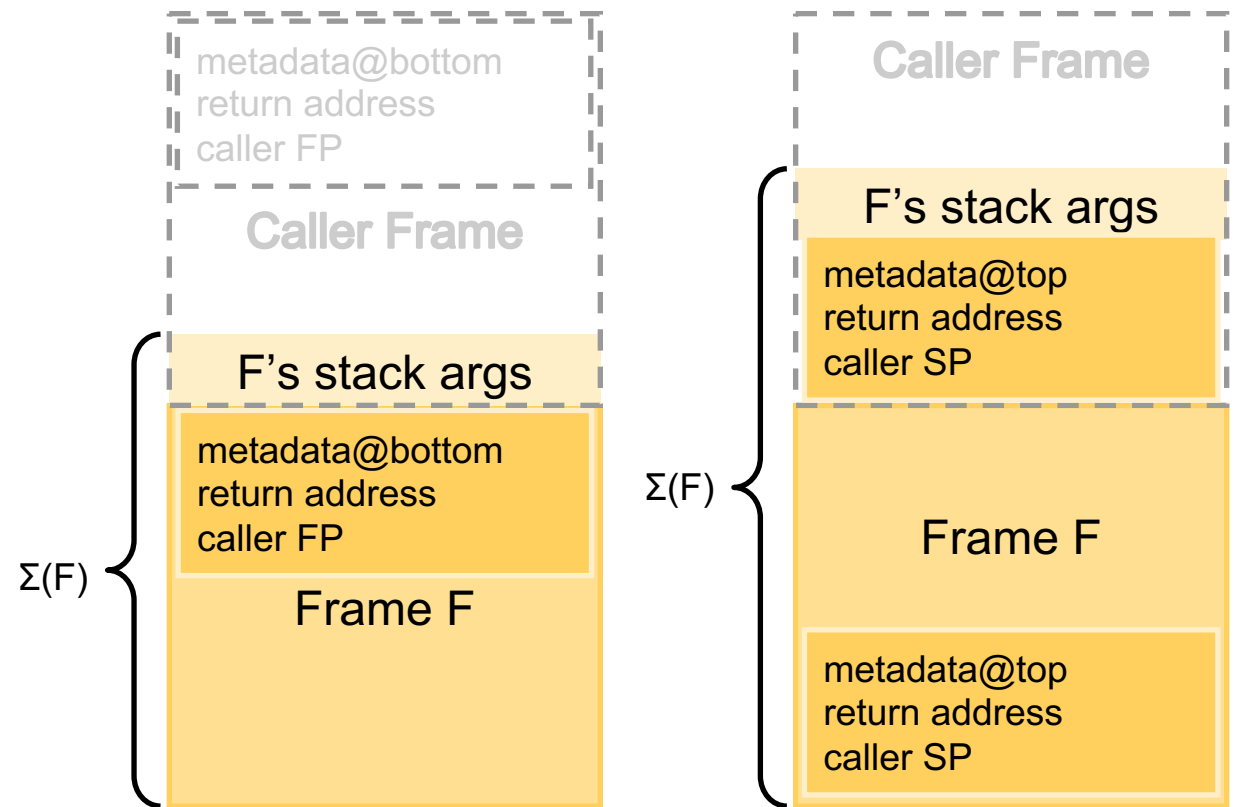
Metadata: return address, prev. FP, prev. SP, etc.

x86, aarch64

- metadata is stored at frame bottom (metadata@bottom)
- all metadata received from the caller is stored in own metadata

ppc64

- metadata at frame top
- not all metadata received from the caller is stored in own metadata
- return address is stored in caller's metadata
- Back chain: caller's SP is stored in own metadata, such that  
`callers_SP = *(calles_SP)`

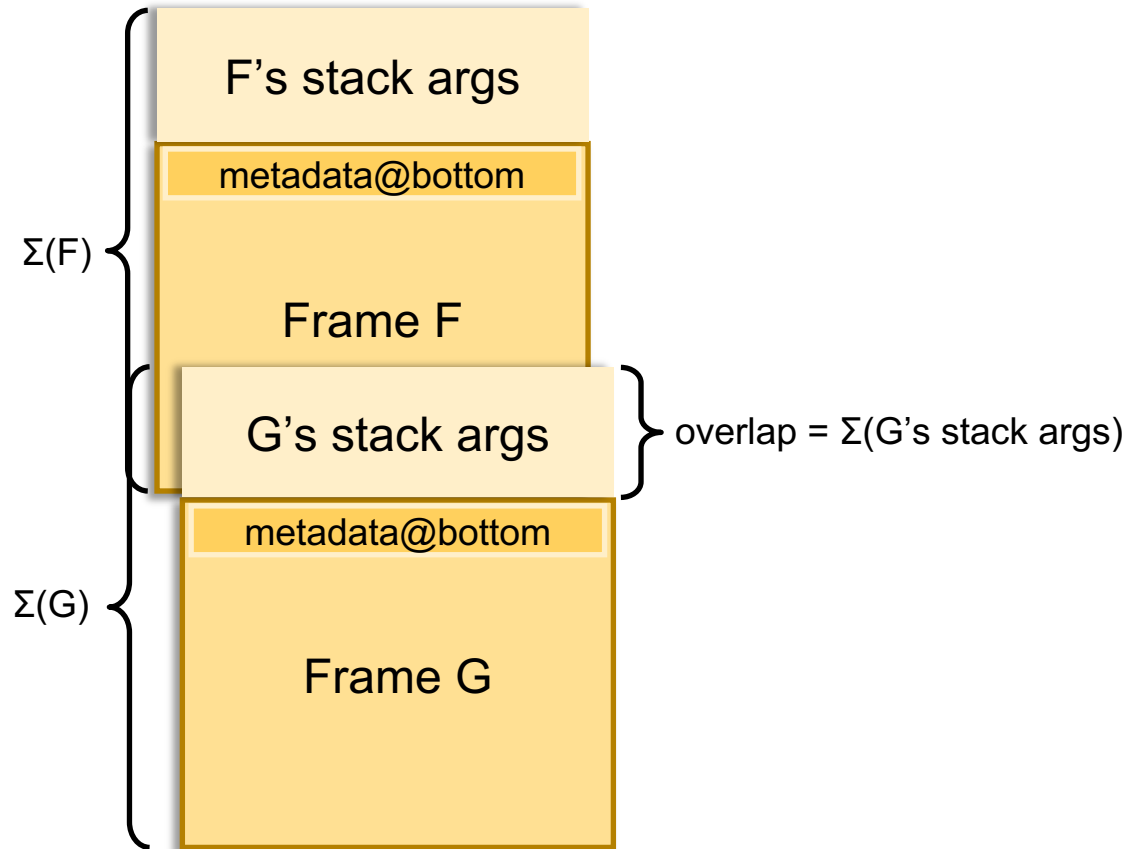


$$\Sigma(F) = \text{frame\_size}(F) + \Sigma(\text{F's stack args}) + \Sigma(\text{metadata@top})$$

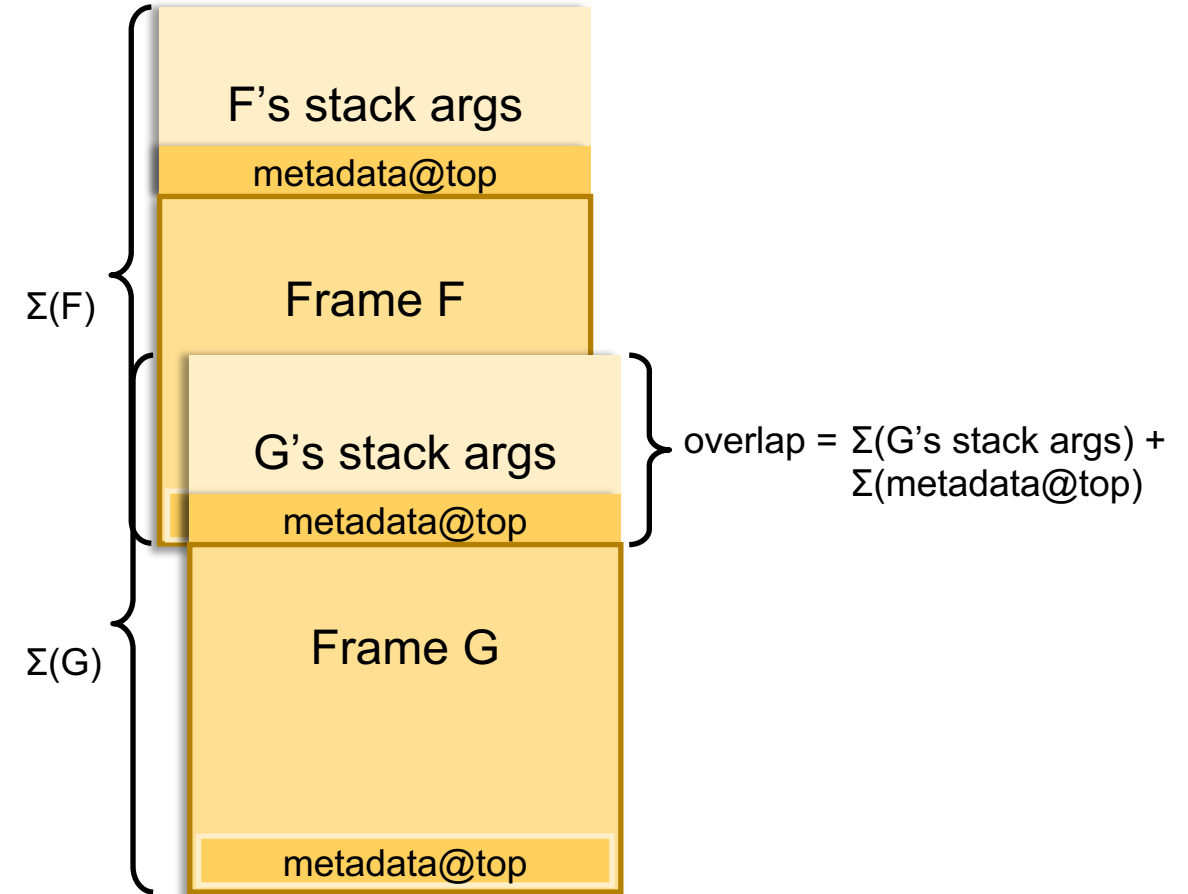
$\Sigma(\text{metadata@top})$  is a platform dependent constant

# Freezing Frames – Slow Path – Sizing Caller/Callee – Metadata

Same Frame Type, metadata@bottom



Same Frame Type, metadata@top

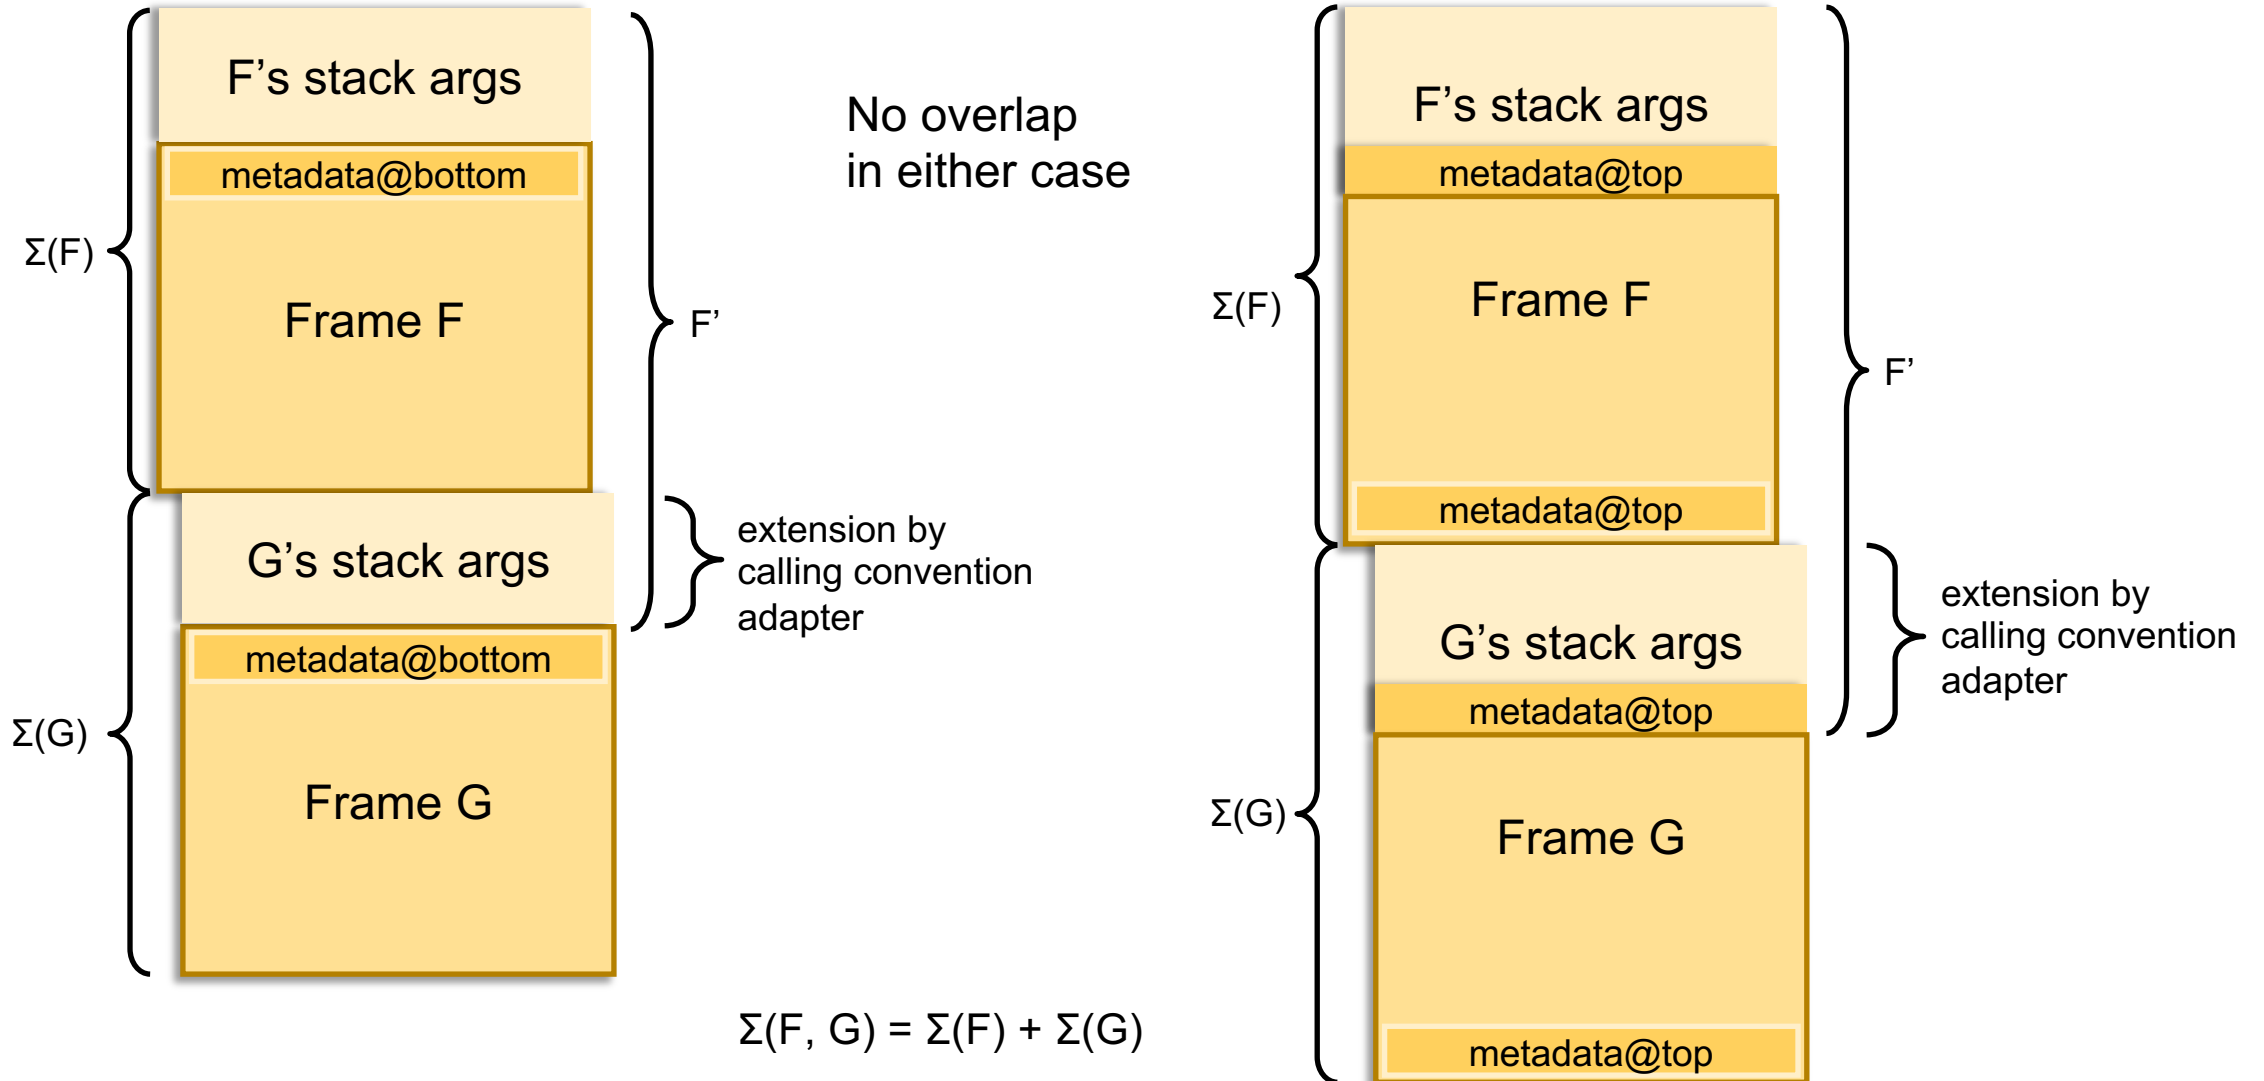


$$\Sigma(F, G) = \Sigma(F) + \Sigma(G) - \text{overlap}$$

# Freezing Frames – Slow Path – Sizing Caller/Callee – Metadata

Different Frame Types, metadata@bottom

Different Frame Type, metadata@top




# Freezing Frames – Slow Path – Fixups

## Relativization

- Interpreter frames have pointers into own frame: `locals`, `esp`, `monitors(, ...)`
- each of them converted into an FP offset

## Patching

- bottom frame: return barrier (from thaw) gets replaced with real return address, i.e. caller's pc
- interpreter frame: `sender_sp`, `back link` are relativized
- no fixup of back link of compiled frames





# Unextended SP

Frames get extended

- by (Calling Convention) adapters to accommodate stack args as required by the callee's calling convention
- by the interpreter to accommodate locals that are not method parameters

A frame's SP before getting extended is the **unextended SP**

- In adapter and interpreter coding it is called **sender SP**
- passed to callee (ppc64: R21\_sender\_SP, x86\_64: r13)
  - ignored if compiled
  - stored in interpreter state if interpreted
- Accessible for runtime code through **frame::unextended\_sp()**  
Not always accurate, e.g., `frame::sender_for_compiled_frame()` sets `unextended_SP = SP` (which is correct if the caller is also compiled and irrelevant (not used), if it is interpreted)

Usage

- address values in a compiled frame
- restore original frame (actually `sender_sp` is used for that)

Hardly uses in shared code before Loom

# Heap Frame Properties

Many uses of `frame::unextended_sp()` in continuations shared code

- Expectation: stack arguments are located @ unextended SP
- tweaked unextended SP of interpreted heap frames on PPC64
  - stack arguments @ unextended SP +  $\Sigma(\text{metadata@top})$
  - with that tweak interpreted and compiled frames on heap are uniform
  - see `ContinuationHelper::InterpretedFrame::patch_sender_sp()`
- Changed shared code to expect stack arguments @ unextended SP +  $\Sigma(\text{metadata@top})$ 
  - Introduced platform dependent constant `frame::metadata_words_at_top` for that

No alignment

- Slow path freeze even removes alignment padding

No backlink for compiled frames

- redundant as it can be computed from the SP + `frame_size`
- PPC64: back chain gets reconstructed when thawing

# Freezing Frames – Fast Path

## Fast Freeze

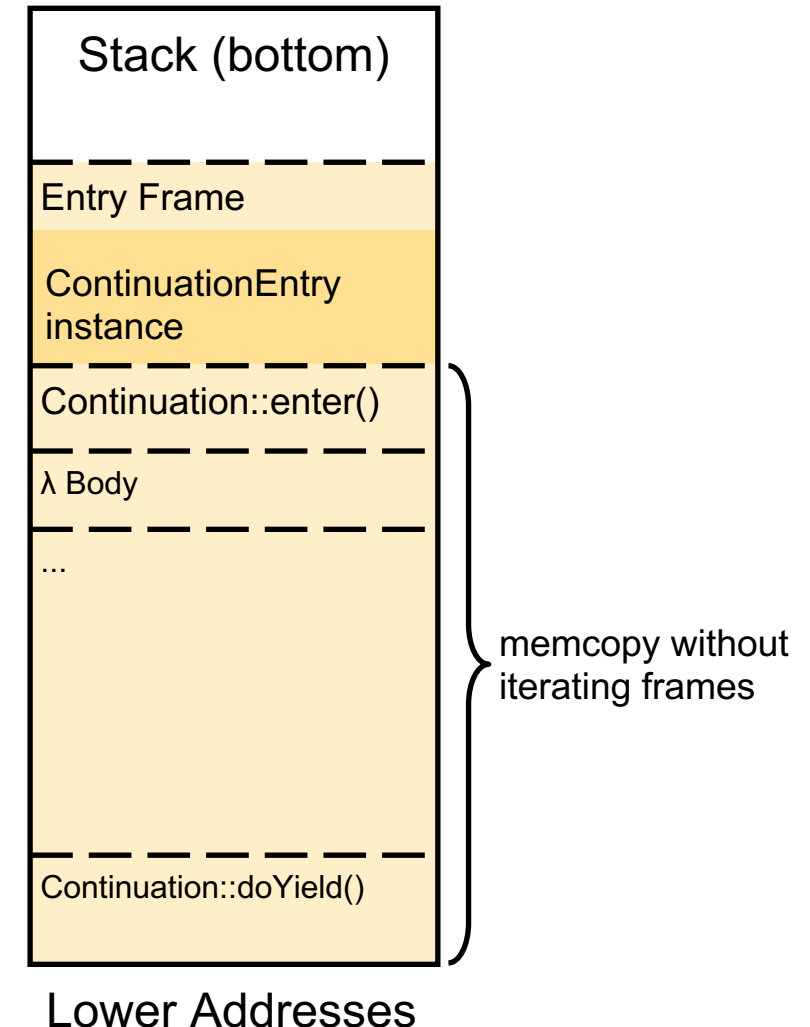
- without iteration of stack frames
- memcpy all between entry and top frame

## Requirements

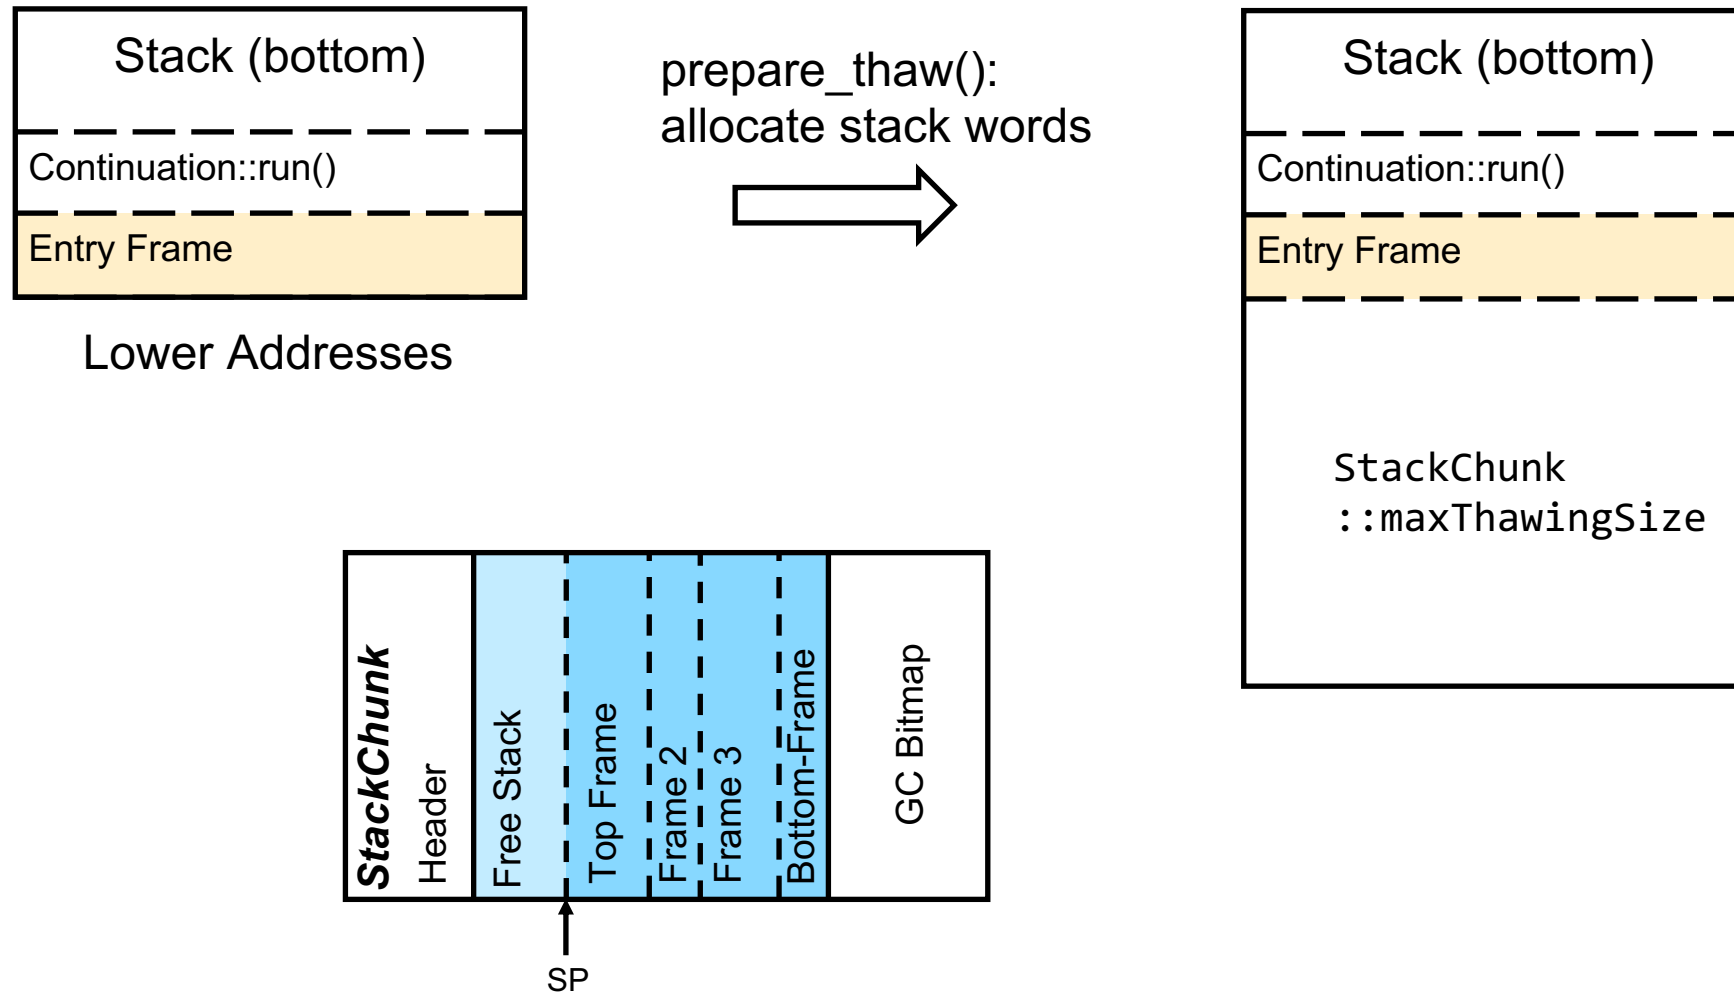
- no locks: check `JavaThread::_held_monitor_count`
- no interpreted / native frames
  - check `JavaThread::_cont_fastpath`: SP of oldest interp./native frame known
  - updated at i2c transitions / deoptimization
  - if interpreted or native frame exists then a i2c transition is needed to freeze
    - `doYield()` uses compiled calling convention
    - JNI to java calls use interpreted calling convention

## No Fixups

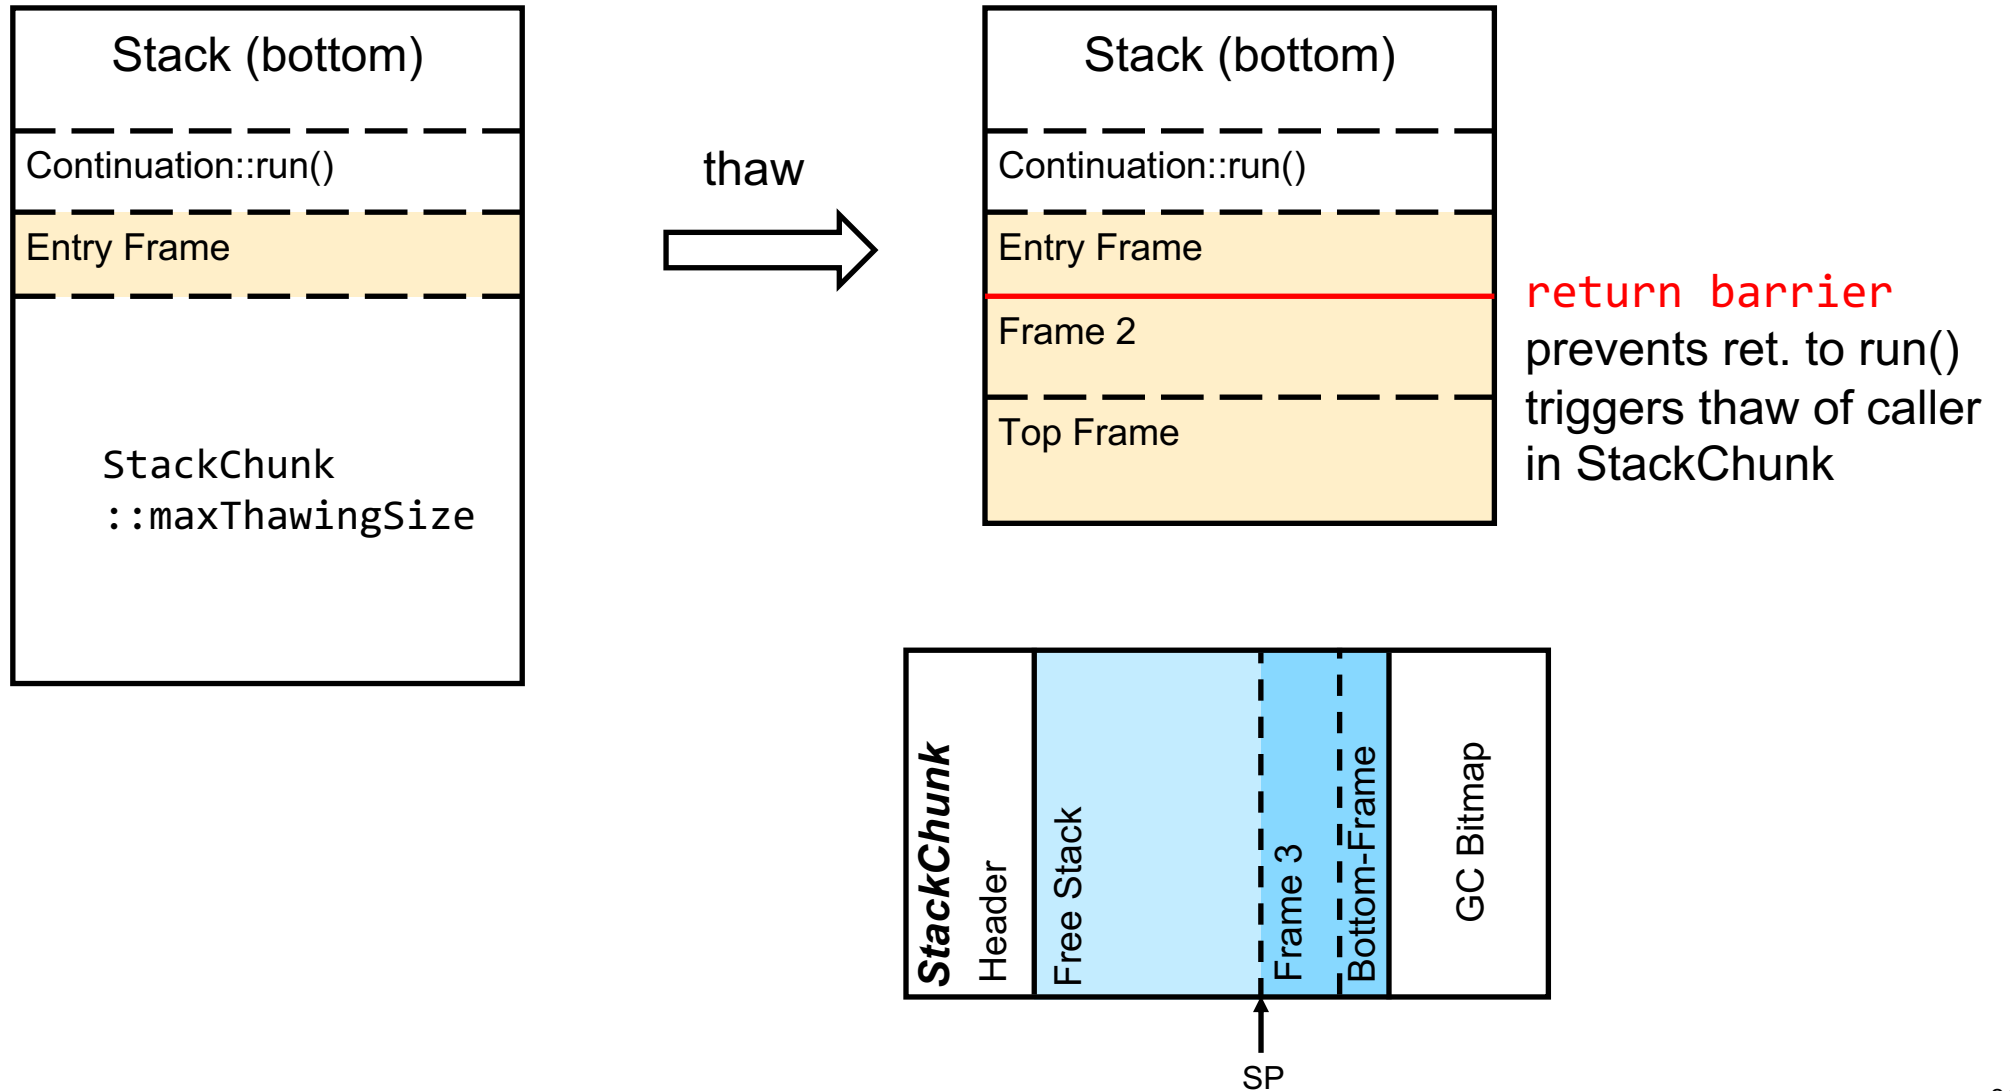
- no stack addresses in compiled frames
- except on ppc: back chain, redundant on heap



## Thawing frames – Slow Path – Diagram (1/2)



## Thawing frames – Slow Path – Diagram (2/2)



# Thawing frames – Preparation

Injected field `maxThawingSize` of `StackChunk C`

- size needed to thaw all frames in `C` with correct alignment
- used to allocate words on stack for frames to be thawed
- actual thaw happens in runtime

# Thawing frames – Slow Path

Recursively iterate frames in StackChunk C on heap

- effectively 1 - 3 frames are thawed
  - limit: `int num_frames = (return_barrier ? 1 : 2);`  
!return\_barrier: first thaw for run() call; 2 topmost frames are known to return
  - one extra frame can be thawed to avoid situations where the top frame in C is compiled and its callee on stack is interpreted ("it makes detecting that situation and adjusting unextended\_sp tricky")
- code: `ThawBase::new_stack_frame()` and its callers

## Fixups

- derelativize interpreter frames
- set correct back link of compiled frames
- uncompress oops, derelativize derived pointers
- set return\_barrier as return address for bottom frame iff C is not empty after thaw to thaw the bottom frames caller

# Thawing frames – Fast Path

## Fast thaw

- without iteration of heap frames
- memcpy all in StackChunk

## Requirements

- only compiled frames in StackChunk
- not transformed by GC (see GC Integration)
- no special GC barriers needed for StackChunk (Shenandoah, ZGC)
- !PreserveFramePointer because otherwise the frames need to be iterated to setup the FramePointers
- Size of frames in StackChunk < 500 words

## No Fixups

- Exception: back chain needs to be restored on ppc64  
(see `Thaw<ConfigT>::patch_caller_links()`)



# GC Integration

new Klass: InstanceStackChunkKlass

- Used to get size of a StackChunk (instances differ in sizes)
- find references in stack using the bitmap stored in each StackChunk instance

StackChunks are lazily transformed for GC just before being traversed

- initialization of oop bitmap
- compressing of oops
- relativization of derived pointers, i.e. converted to offsets  
Derived pointers are pointers into heap objects
- after transformation frames cannot be frozen or thawed on the fast path
  - freeze: difficult to synchronize the necessary transformation with concurrent GCs
  - thaw: would require iterating the frames to uncompress oops and derelativize derived pointers
- see `stackChunkOopDesc::transform()`

Bitmap is cleared when thawing

- only the part that covers the stack args of the bottom frame being thawed
- no need to clear bitmap for all thawed frames because the bits covering free region are ignored

Transformed StackChunks cannot accommodate new frames anymore

# Potential Performance Cliffs

If a StackChunk is transformed for GC

- it cannot accommodate new frames  
a new one must be allocated if the continuation yields again
- its frames cannot be fast thawed

Large StackChunks (> 500 words) cannot be fast thawed

Not compilable methods prevent fast path context switch

# Integration with established Thread and Stack related (internal) VM APIs

## Iterating Physical Frames (class frame)

- RegisterMap: holds stack walk state, e.g. to find callee saved registers
  - `_chunk`: StackChunk currently walked
  - `_walk_cont`: where to continue a walk at the callee of the entry frame
    - either with the top frame on heap, used, e.g., in JVMTI
    - or with the entry frame
- `frame::_on_heap` indicates if a frame resides in a StackChunk on heap
- `frame::sender_raw()` uses `stackChunkOopDesc::sender(frame&)` to get the sender of a heap frame
- StackChunkFrameStream
  - Iterator for frames in a StackChunk
  - used by `stackChunkOopDesc::sender(frame&)`
- frame instances are **relativized** if representing heap frames
  - heap pointers (`_sp`, `_fp`) are offsets relative to the StackChunk (wrapped in handle)
  - relativized frames are not invalidated by safepoints
  - derelativize before use, e.g. `stackChunkOopDesc::interpreter_frame_method()`

# Integration with established Thread and Stack related (internal) VM APIs

Virtual Frames (class `vframe`)

- abstracts away details like relativized frames

JavaThread

- `_vthread`: returned by `j.l.Thread.currentThread()`
- `_cont_entry`: list head of `ContinuationEntry`s
- `vthread_continuation()`: finds first `VirtualThread` continuation in `_cont_entry` list.

# Testing

`make test TEST="hotspot_loom jdk_loom" TEST_VM_OPTS="-XX:+VerifyContinuations"`

- x86\_64, aarch64, ppc64le
- fastdebug, release

# References

JBS Item: <https://bugs.openjdk.org/browse/JDK-8286302>

Code: [https://github.com/reinrich/jdk/tree/ppc\\_port](https://github.com/reinrich/jdk/tree/ppc_port)

# Thank you.

Contact information:

Richard Reingruber  
richard.reingruber@sap.com